

Chapter 6: Comparing Two Populations

Introductory Statistics for Engineering Experimentation

Peter R. Nelson, Marie Coffin and Karen A.F. Copeland

Slides by Douglas Bates

Outline

6.1 Paired samples

6.2 Independent samples

6.3 Comparing two binomial proportions

Outline

6.1 Paired samples

6.2 Independent samples

6.3 Comparing two binomial proportions

Outline

6.1 Paired samples

6.2 Independent samples

6.3 Comparing two binomial proportions

Comparative experiments

- The “single sample” types of hypothesis tests described in the last chapter, where $H_0 : \mu = \mu_0$ or $H_0 : p = p_0$, are often used to compare a modified process to a standard method.
- Even if we reject H_0 in favor of H_a all we can really conclude is that the process has changed. We don't know if the change is due to our modification (the experimental factor) or due to other environmental factors.
- If we wish to focus on a particular modification it is better to use a *comparative experiment* in which we keep environmental factors (raw materials, time, temperature, etc.) as consistent as possible and change only the factor of interest.
- In this chapter we focus on comparative experiments where the experimental factor has only two levels. In statistical terms we are comparing two populations, corresponding to the two levels of the factor.

Overview of techniques

- When the data are on a continuous scale, we wish to compare the means, written μ_1 and μ_2 of the two populations, either by forming a confidence interval on $\mu_1 - \mu_2$ or testing the hypothesis $H_0 : \mu_1 = \mu_2$ versus a one- or two-sided alternative.
- If we have controlled for a known source of variability by taking, say, before-after measurements on the same subject, we consider the data as a set of n pairs, $(y_{1i}, y_{2i}), i = 1, \dots, n$ and analyze the differences $d_i = y_{1i} - y_{2i}$ as a single sample.
- For unpaired data we take the difference in the sample means, suitably standardized, and compare to a T distribution in which we approximate the degrees of freedom.
- For binary response data we compare the observed proportions \hat{p}_1 and \hat{p}_2 using a standardized statistic.

R functions used in this chapter

- Comparison of two samples on a continuous scale is done with `t.test`, as in the previous chapter. Paired samples are indicated by the optional argument `paired = TRUE`.
- We can always use the data in the “stacked” format where all the response measurements are in one column and there is a second column, a factor with two levels, that distinguishes the two samples.
- When the sample sizes are equal, $n_1 = n_2$, and especially for paired samples, the data are often available in an “unstacked” format. That is, the responses are in two columns.
- The `t.test` function is used in both cases but the form of the arguments is different. For stacked data we can use a formula/data specification.
- Comparison of two population proportions is done with `prop.test`. All we need for this test are the sample sizes, n_1 and n_2 , and the number of successes in each sample, y_1 and y_2 .

Section 6.1, Paired samples

- The trick with paired samples is recognizing that the observations in the two samples are paired.
- Obviously, if they are to be paired you must have equal sample sizes, $n_1 = n_2$.
- There must also be some other factor (subject, location, raw material, etc.) that associates the first observation in sample 1 with the first observation in sample 2, and so on.
- A scatterplot of y_{2i} versus y_{1i} should show points scattered about a line with positive slope. If it doesn't then the pairing is unsuccessful.
- To analyze the data we take the differences, $d_i = y_{1i} - y_{2i}, i = 1, \dots, n$ and analyze them as a single sample. The hypothesis $H_0 : \mu_1 = \mu_2$ corresponds to $H_0 : \mu_d = 0$.
- In practice we can specify `paired = TRUE` in the call to `t.test` to have the data analyzed as a paired sample.

Example 6.1.1

- The `uvcoatin` data are from an experiment comparing the two UV coatings on lenses. Each pair of observations came from one pair of glasses worn by a person for 3 months, one lens with the new coating and one lens with the commercial coating.

```
> str(uvcoatin)
```

```
'data.frame': 10 obs. of 3 variables:
```

```
$ a : num 8.9 9.4 11.2 11.4 13 6.4 13.4 5.6 4.8 15.8
```

```
$ b : num 8.5 9.3 10.8 11.6 12.9 6.5 13.1 5.1 4.3 15.6
```




```
$ diff: num 0.4 0.1 0.4 -0.2 0.1 -0.1 0.3 0.5 0.5 0.2
```

```
> head(uvcoatin)
```

	a	b	diff
1	8.9	8.5	0.4
2	9.4	9.3	0.1
3	11.2	10.8	0.4
4	11.4	11.6	-0.2
5	13.0	12.9	0.1
6	6.4	6.5	-0.1

Plots of UV coating data

- We can plot these data as a scatterplot. Sometimes it is helpful also to rotate the scatterplot by 45° , which corresponds to plotting the difference versus the mean. The `tmd` function automates this.
- We should also check a normal probability plot of the differences.



Paired t-test on UV coating data

```
> with(uvcoatin, t.test(a, b, alt = "g", paired = TRUE))
```

```
Paired t-test
```

```
data: a and b
```

```
t = 2.8508, df = 9, p-value = 0.009533
```

```
alternative hypothesis: true difference in means is greater than 0
```

```
95 percent confidence interval:
```

```
0.07853456      Inf
```

```
sample estimates:
```

```
mean of the differences
```

```
0.22
```

```
> with(uvcoatin, t.test(a, b, pair = 1))$conf.int
```

```
[1] 0.04542425 0.39457575
```

```
attr(,"conf.level")
```

```
[1] 0.95
```

Section 6.2, Independent samples

- If we have independent (i.e. unpaired) samples, of sizes n_1 and n_2 , from two populations that we assume have $\mathcal{N}(\mu_1, \sigma_1^2)$ and $\mathcal{N}(\mu_2, \sigma_2^2)$ distributions, the test statistic for $H_0 : \mu_1 = \mu_2$ versus a one- or two-sided alternative is

$$t_{\text{obs}} = \frac{\bar{y}_{1\cdot} - \bar{y}_{2\cdot}}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

- In theory, this does not have exactly a T distribution. In practice, it is close enough but the effective degrees of freedom, ν , depend on the relative variability of the two samples. If $s_1^2 \approx s_2^2$ then $\nu \approx (n_1 - 1) + (n_2 - 1)$. If $s_1^2 \gg s_2^2$ then $\nu \approx n_1 - 1$ and vice versa. We use

$$\nu = \frac{\left(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}\right)^2}{\frac{1}{n_1-1} \left(\frac{s_1^2}{n_1}\right)^2 + \frac{1}{n_2-1} \left(\frac{s_2^2}{n_2}\right)^2}$$


Example 6.2.1

```
> str(railcar3)
```

```
'data.frame': 17 obs. of 2 variables:
```

```
$ type      : Factor w/ 2 levels "A","B": 1 1 2 1 2 2 2 1 1 2 ...
```

```
$ moisture: num  3.4 3.1 2.3 2.6 2.1 2.8 2.5 3 3.2 2.6 ...
```



```
> t.test(moisture ~ type, railcar3, alt = "g")
```

```
Welch Two Sample t-test
```

```
data: moisture by type
```

```
t = 2.8646, df = 13.861, p-value = 0.006294
```

```
alternative hypothesis: true difference in means is greater than 0
```

```
95 percent confidence interval:
```

```
0.1549536      Inf
```


```
sample estimates:
```

```
mean in group A mean in group B
```

```
2.925000
```

```
2.522222
```

Examples 6.2.4 and 6.2.5



```
> t.test(yield ~ process, assay)
```

Welch Two Sample t-test

data: yield by process

t = -8.0993, df = 4.901, p-value = 0.000511

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-53.15548 -27.42452

sample estimates:

mean in group A mean in group B

39.45

79.74

Section 6.3, Comparing two binomial proportions

- Our data are y_1 and y_2 , the numbers of successes, and n_1 and n_2 , the numbers of trials. We consider difference in the observed proportions, $\hat{p}_1 - \hat{p}_2$, with a standard error of

$$\sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}$$

- The calculation done in the `prop.test` function in R uses what is called a “continuity correction” and is slightly more accurate than the one described in the text. For the data in Example 6.3.1 it provides

```
> prop.test(c(8, 14), c(500, 500))$conf.int
```

```
[1] -0.032167502  0.008167502
```

```
attr(,"conf.level")
```

```
[1] 0.95
```

Hypothesis tests

- We can also use `prop.test` for hypothesis tests. For example 6.3.4 the results are

```
> prop.test(c(26, 17), c(312, 329))
```

```
2-sample test for equality of proportions with  
continuity correction
```

```
data:  c(26, 17) out of c(312, 329)
```

```
X-squared = 2.0841, df = 1, p-value = 0.1488
```

```
alternative hypothesis: two.sided
```

```
95 percent confidence interval:
```

```
-0.01035392  0.07367712
```

```
sample estimates:
```

```
prop 1      prop 2
```

```
0.08333333 0.05167173
```

- The test statistic quoted here is the square of the z statistic described in the text. (The name "X-squared" stands for χ^2 . The distribution of the square of a standard normal is a χ^2 on 1 degree of freedom.)

Turning off the continuity correction

- The results using the continuity correction are preferred but some people insist on the less accurate version because it corresponds to the text book formula. Use the optional argument `correct = FALSE`

```
> sqrt(prop.test(c(26, 17), c(312, 329))$statistic) # with cont.
```

```
X-squared  
1.443652
```

```
> sqrt(prop.test(c(26, 17), c(312, 329), corr = 0)$statistic)
```

```
X-squared  
1.601594
```

Sample sizes and power

- The function `power.prop.test` provides sample size calculations or power calculations for tests of proportions.
 - For sample size calculations we need working values of the two proportions plus the significance level, α , and the power, $1 - \beta$. If only the deviation in the proportions is specified, use the “worst case scenario” of equal spacing around $p_1 = 0.5$.
 - Example 6.3.5 asks for sample sizes when $\alpha = 0.1$ and the power should be 0.9 when $p_A - p_B = 0.1$.
- > `power.prop.test(p1 = 0.45, p2 = 0.55, sig = 0.1, power = 0.9,`

```
Two-sample comparison of proportions power calculation
  n = 326.8305
  p1 = 0.45
  p2 = 0.55
sig.level = 0.1
  power = 0.9
alternative = one.sided
```

NOTE: n is number in *each* group