

Chapter 9: Multi-factor experiments

Introductory Statistics for Engineering Experimentation

Peter R. Nelson, Marie Coffin and Karen A.F. Copeland

Slides by Douglas Bates

Outline

9.1 ANOVA for multi-factor experiments

9.2 2^k factorial designs

Outline

9.1 ANOVA for multi-factor experiments

9.2 2^k factorial designs

Chapter 9: Multi-factor experiments

- We are not limited to using two factors in an experiment. Especially in screening experiments, where we are seeking to distinguish the “important few” factors from the “trivial many”, it is useful to incorporate many factors.
- Incorporating many factors, each with several levels, results in a large number of combinations of levels. If we want replicates the number of observations required is even larger.
- One way to reduce the number of observations required is to use a small number of levels. We often use just two levels.
- At two levels, there is no distinction between a numeric covariate, a categorical covariate or an ordered category. For technical reasons we often represent such a factor as an ordered category.

Section 9.1: ANOVA for multi-factor experiments

- For a replicated design we proceed as for two-factor experiments and fit a model with all possible interactions.
- We check the highest-order interaction first. If it is not significant we re-fit without that term and continue checking. In a balanced factorial design omitting certain interactions will not change the estimates of coefficients for other interactions, but it will change MS_e .
- The formula for model with all possible interactions has * between the factors.
- We can modify the formula by “subtracting” particular terms using the `update` function.

Example 9.1.1: Battery separator data

```
> str(separate)
```


```
'data.frame': 24 obs. of 4 variables:
```

```
$ silica: Ord.factor w/ 2 levels "Low"<"High": 1 1 1 1 1 1 1 1 1 1 1 ...
```

```
$ time : Ord.factor w/ 2 levels "Low"<"High": 1 1 1 1 1 1 1 2 2 2 2 ...
```

```
$ temp : Ord.factor w/ 2 levels "Low"<"High": 1 1 1 2 2 2 1 1 1 2 ...
```

```
$ y : num 40.9 42.2 41.3 46.3 47 48.2 48.6 49.5 46.6 69 ...
```



Example 9.1.1 (cont'd)

```
> summary(fm1 <- aov(y ~ time * temp * silica, separate))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
time	1	1732.30	1732.30	988.239	8.212e-16
temp	1	1318.68	1318.68	752.280	7.014e-15
silica	1	129.27	129.27	73.746	2.185e-07
time:temp	1	31.97	31.97	18.238	0.0005854
time:silica	1	78.84	78.84	44.979	5.041e-06
temp:silica	1	38.25	38.25	21.823	0.0002553
time:temp:silica	1	110.51	110.51	63.044	6.113e-07
Residuals	16	28.05	1.75		

- We see that the highest-order interaction term is significant. In such cases we generally retain all the lower-level interactions and the main effects whether or not they have small p-values. (In this case they all do.)

Example 9.1.1 (cont'd)

- There are two approaches for follow-up. We can consider the 8 different combinations of factor levels as 8 levels in a one-factor model or we can condition on a level of one factor and consider the others.

```
> summary(fm1a <- aov(y ~ time * silica, separate,
+                      subset = temp == "High"))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
time	1	1117.47	1117.47	567.2437	1.029e-08
silica	1	154.08	154.08	78.2149	2.107e-05
time:silica	1	1.33	1.33	0.6768	0.4345
Residuals	8	15.76	1.97		

```
> summary(fm1b <- aov(y ~ time + silica, separate,
+                      subset = temp == "High"))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
time	1	1117.47	1117.47	588.371	1.646e-09
silica	1	154.08	154.08	81.128	8.483e-06
Residuals	9	17.09	1.90		

Blocking with multiple experimental factors

- It is not uncommon to have blocking factors in a multi-factor design. These are known sources of variability that cannot (or should not) be held constant.
- As before, we list blocking factors first in any model formula. This is not important for balanced data but can be important with unbalanced data.

```
> str(defect <-
+   data.frame(num = c(9,9,5,5,7,8,2,3,8,7,4,4),
+                 op = gl(3,4),
+                 app = gl(2,2,12,labels = c("A","B")),
+                 size = gl(2,1,12)))
```

```
'data.frame': 12 obs. of 4 variables:
 $ num : num 9 9 5 5 7 8 2 3 8 7 4 4
 $ op : Factor w/ 3 levels "1","2","3": 1 1 1 1 2 2 2 2 3 3 ...
 $ app : Factor w/ 2 levels "A","B": 1 1 2 2 1 1 2 2 1 1 ...
 $ size: Factor w/ 2 levels "1","2": 1 2 1 2 1 2 1 2 1 2 ...
```

Structure of defect data

```
> ftable(xtabs(num ~ op+size+app, defect)) # compare Table 9.7
```

```
      app A B
op size
1  1      9 5
   2      9 5
2  1      7 2
   2      8 3
3  1      8 4
   2      7 4
```

```
> summary(fm2 <- aov(num ~ op + size * app, defect))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
op	2	8.167	4.083	9.8	0.01287
size	1	0.083	0.083	0.2	0.67041
app	1	52.083	52.083	125.0	3.056e-05
size:app	1	0.083	0.083	0.2	0.67041
Residuals	6	2.500	0.417		

Reduced models

```
> summary(fm3 <- aov(num ~ op + size + app, defect))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
op	2	8.167	4.083	11.0645	0.006803
size	1	0.083	0.083	0.2258	0.649120
app	1	52.083	52.083	141.1290	6.803e-06
Residuals	7	2.583	0.369		

```
> summary(fm4 <- aov(num ~ op + app, defect))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
op	2	8.167	4.083	12.25	0.003671
app	1	52.083	52.083	156.25	1.570e-06
Residuals	8	2.667	0.333		

- There is no need for multiple comparisons of levels of **app** (there are only two levels).
- Generally we don't compare levels of **op** because it is a blocking factor. In this case we may want to investigate why there is such a large variability between operators.

Unreplicated multi-factor designs

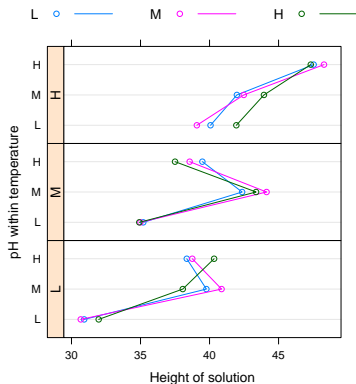
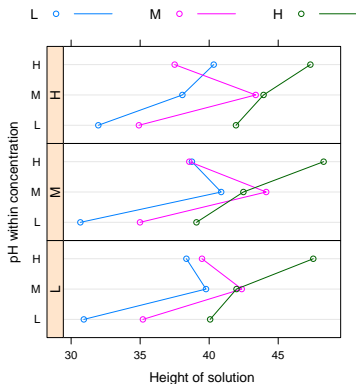
- If we allow the highest-order interaction term in an unreplicated factorial design we will have no degrees of freedom for error.
- We can remove the highest-order interaction by using “powers” in the formula to restrict to interactions up to a specific order.

```
> str(defoam)
```

```
'data.frame': 27 obs. of 4 variables:
```

```
$ conc : Ord.factor w/ 3 levels "L"<"M"<"H": 1 1 1 1 1 1 1 1 1 2 ...  
$ pH    : Ord.factor w/ 3 levels "L"<"M"<"H": 1 1 1 2 2 2 3 3 3 1 ...  
$ temp  : Ord.factor w/ 3 levels "L"<"M"<"H": 1 2 3 1 2 3 1 2 3 1 ...  
$ height: num 30.9 35.2 40.1 39.8 42.4 ...
```

Plots of the defoam data



- Note that there are effectively no differences due to concentration, at the levels examined.

Models for the defoam data

```
> summary(fm5 <- aov(height ~ (conc + pH + temp)^2, defoam))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
conc	2	0.759	0.380	0.3108	0.7413502
pH	2	238.977	119.488	97.7850	2.385e-06
temp	2	228.653	114.327	93.5608	2.826e-06
conc:pH	4	4.058	1.014	0.8301	0.5420108
conc:temp	4	2.949	0.737	0.6034	0.6712171
pH:temp	4	82.012	20.503	16.7789	0.0005885
Residuals	8	9.776	1.222		

```
> summary(fm6 <- aov(height ~ (pH + temp)^2, defoam))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
pH	2	238.977	119.488	122.609	3.270e-11
temp	2	228.653	114.327	117.312	4.733e-11
pH:temp	4	82.012	20.503	21.038	1.378e-06
Residuals	18	17.542	0.975		

Section 9.2: 2^k factorial designs

- A commonly used type of screening design is a 2^k factorial design in which each of k factors is examined at 2 levels only.
- Such a design will only detect linear trends and possible interactions. It cannot, for example, detect a quadratic relationship reliably (Fig. 9.9).
- One technical point, it can be an advantage to represent the 2-level factors as `ordered` and to use the “Helmert” contrasts for ordered factors. These provide a $+/-$ coding in the model matrix.

```
> options(contrasts = c("contr.treatment", "contr.helmert"))
> str(dents <- data.frame(ok = c(917,600,953,750,735,567,977,647)
+                           A = gl(2,1,8,ord=1), B = gl(2,2,8,ord=
+                           C = gl(2,4,ord=1)))
```

```
'data.frame': 8 obs. of 4 variables:
```

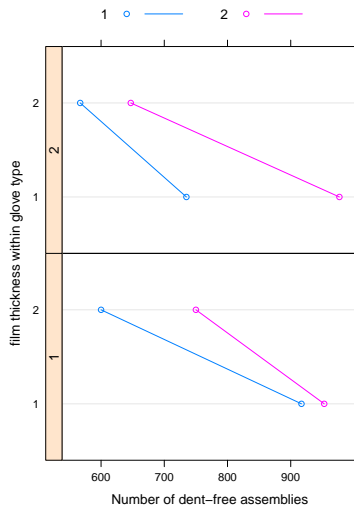
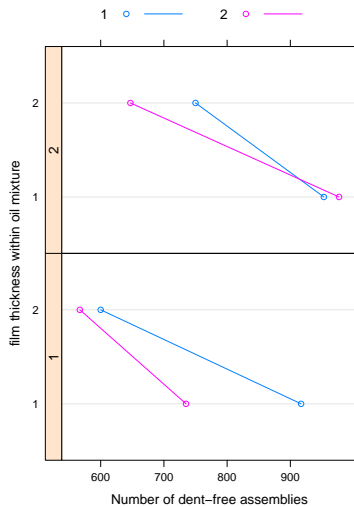
```
$ ok: num 917 600 953 750 735 567 977 647
```

```
$ A : Ord.factor w/ 2 levels "1"<"2": 1 2 1 2 1 2 1 2
```

```
$ B : Ord.factor w/ 2 levels "1"<"2": 1 1 2 2 1 1 2 2
```

```
$ C : Ord.factor w/ 2 levels "1"<"2": 1 1 1 1 2 2 2 2
```

Plots of the dents data



Model fits for the dents data

```
> summary(fm7 <- aov(ok ~ (A+B+C)^2, dents))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
A	1	129540	129540	13.6043	0.1685
B	1	32258	32258	3.3877	0.3168
C	1	10804	10804	1.1347	0.4799
A:B	1	288	288	0.0302	0.8904
A:C	1	60	60	0.0064	0.9494
B:C	1	2312	2312	0.2428	0.7085
Residuals	1	9522	9522		

None of the two-factor interactions are significant (but it is very difficult to show significance with 1 degree of freedom for residuals).

```
> summary(fm8 <- aov(ok ~ A+B+C, dents))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
A	1	129540	129540	42.5333	0.002854
B	1	32258	32258	10.5916	0.031243
C	1	10804	10804	3.5475	0.132744
Residuals	4	12182	3046		

Model fits for the dents data (cont'd)

```
> summary(fm9 <- aov(ok ~ A+B, dents))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
A	1	129540	129540	28.1769	0.003171
B	1	32258	32258	7.0166	0.045488
Residuals	5	22987	4597		

```
> summary.lm(fm9)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	768.25	23.97	32.047	5.56e-07
A1	-127.25	23.97	-5.308	0.00317
B1	63.50	23.97	2.649	0.04549

Residual standard error: 67.8 on 5 degrees of freedom

Multiple R-squared: 0.8756, Adjusted R-squared: 0.8258

F-statistic: 17.6 on 2 and 5 DF, p-value: 0.005458

```
> fitted(fm9)
```

1	2	3	4	5	6	7	8
832.0	577.5	959.0	704.5	832.0	577.5	959.0	704.5


Interpretation of coefficients

- The coefficients shown in the `summary.lm` output are in the $+/-$ encoding. The `(Intercept)` is the response for a “typical” setting; the `A1` coefficient is added to this for the high level of `A` and subtracted for the low level. Similarly for `B1`.
- That is, the fitted values are the `(Intercept)` plus the four possible combinations of $\pm A1$ and $\pm B1$.
- *Yates' algorithm* is a method for determining these coefficient estimates by hand calculation. It is interesting but no longer needed.
- Notice also that the standard errors of all the coefficients are identical. This provides an alternative approach to evaluating the significance of coefficients in a *saturated* model (as many coefficients as there are responses). We examine a normal probability plot of the coefficient estimates (excluding the intercept) which we expect to have outliers. The non-outliers are the “trivial many”. The outliers are the “important few”

QQ-plot of coefficients from the saturated model

```
> (cc <- coef(fm10 <- lm(ok ~ A * B * C, dents))[-1])
```

A1	B1	C1	A1:B1	A1:C1	B1:C1	A1:B1:C1
-127.25	63.50	-36.75	-6.00	2.75	17.00	-34.50

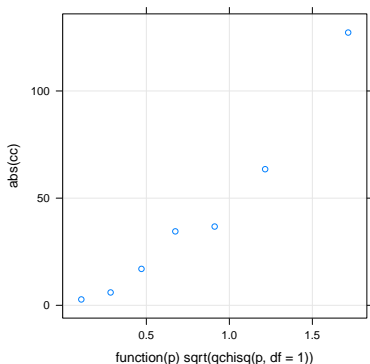


Interpretation of the QQ plot of coefficient estimates

- The plot on the previous slide indicates that 2 coefficients are the “important few”.
- Checking the listing of the coefficients we see that these are the main effects for A and for B — the same conclusion as before.
- Generally this approach is preferred to the approach of testing for two-factor interactions using F tests with only 1 denominator degree of freedom. F tests with very few denominator degrees of freedom are not at all powerful.
- One variation on this approach is to compare the absolute value of the coefficients to the theoretical quantiles. This is called a “half-normal” plot.
- It is not part of the course but the way to create such a plot is shown on the next slide. In this plot you are only interested in the large magnitudes. The “reference line” would be a line passing through the origin. You imagine the low magnitudes as falling near such a line with the high values above the line.

Half-normal plot of saturated model coefficients

```
> print(qqmath( ~ abs(cc), aspect = 1, type = c("g","p"),  
+           distribution = function(p) sqrt(qchisq(p, df = 1))
```



Here you could make a case for 5 values near the line (and 2 above) or for 6 values near the line and only one above.